

# Recurrent Neural Networks

VVIT@2025

November 8, 2025

# Table of Contents

- 1 Introduction
- 2 RNN Architecture
- 3 Equations
- 4 RNN Equations
- 5 Types of RNN Architecture
- 6 Applications
- 7 Introduction to LSTM
- 8 LSTM Architecture
- 9 LSTM Gates
- 10 LSTM Cell State Update
- 11 Applications of LSTM

# Introduction

- Overview of neural networks
- Brief explanation of recurrent neural networks

# What are Recurrent Neural Networks?

- Neural network architecture designed for sequential data processing
- Suitable for tasks involving time series, speech recognition, natural language processing, etc.
- Process input data of arbitrary length and maintain internal state
- Share parameters across different time steps

# Key Concepts

- Hidden State
- Activation Function
- Backpropagation Through Time (BPTT)
- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

# Applications of RNNs

- Language Modeling
- Machine Translation
- Speech Recognition
- Image Captioning
- Sentiment Analysis
- Time Series Prediction

# Training and Architecture

- Training RNNs using Backpropagation Through Time (BPTT)
- Challenges: Vanishing and Exploding Gradients
- Architectural variations: LSTM, GRU
- Deep RNNs and Bidirectional RNNs

# RNN Architecture

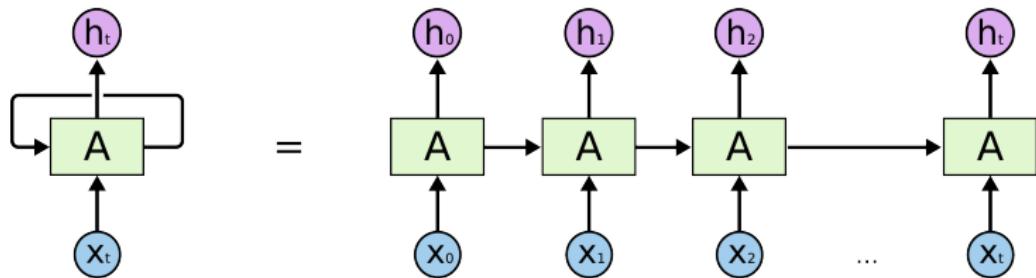


Figure: Illustration of an RNN architecture

# Equations

- RNN forward pass equation:  $h_t = \text{activation}(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$
- RNN backward pass equation:  $\frac{\partial L}{\partial h_t} = \frac{\partial L}{\partial h_{t+1}} \frac{\partial h_{t+1}}{\partial h_t}$

# RNN Equations

- RNN forward pass equation:  $h_t = \text{activation}(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$
- RNN backward pass equation:  $\frac{\partial L}{\partial h_t} = \frac{\partial L}{\partial h_{t+1}} \frac{\partial h_{t+1}}{\partial h_t}$

# Types of RNN Architecture

- Simple RNN (Elman network)
- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

# Applications of RNNs

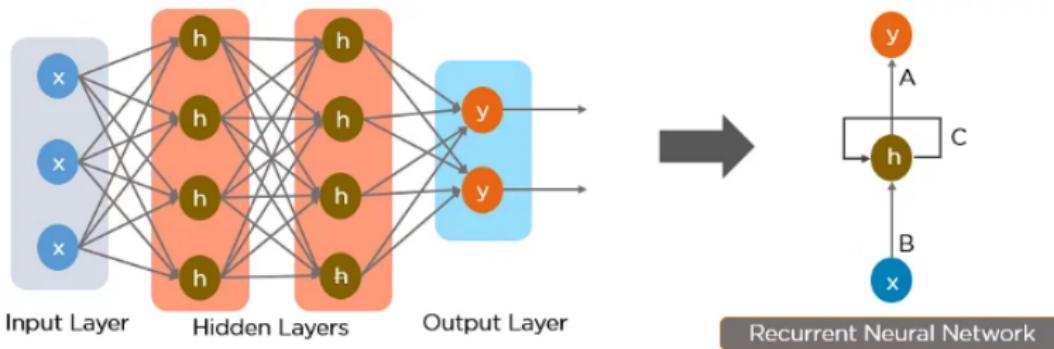
- Natural language processing
- Speech recognition
- Time series analysis

# What Is a Neural Network?

A Neural Network consists of different layers connected to each other, working on the structure and function of a human brain. It learns from huge volumes of data and uses complex algorithms to train a neural net.

# What Is a Recurrent Neural Network (RNN)?

A Recurrent Neural Network (RNN) works on the principle of saving the output of a particular layer and feeding this back to the input in order to predict the output of the layer.



**Figure:** Recurrent Neural Network (RNN) Architecture

The nodes in different layers of the neural network are compressed to form a single layer of recurrent neural networks. A, B, and C are the parameters of the network.

# Introduction

- Humans possess the ability to build upon prior knowledge for continuous understanding.
- Traditional neural networks lack this persistence, hindering sequential reasoning tasks.
- Recurrent neural networks (RNNs) provide a solution with their inherent looping structure.

# Understanding RNNs

- RNNs consist of recurrent connections that allow information to persist across time steps.
- Each time step receives input and produces output, with hidden states acting as memory.
- Hidden states are updated at each step, incorporating current input and previous hidden state.

# The Power of Sequential Reasoning

- Traditional neural networks lack the ability to leverage prior context for sequential tasks.
- RNNs excel in tasks requiring understanding of temporal dependencies and context.
- Example: Movie event classification based on previous events informs later predictions.

# The Power of Sequential Reasoning

In sequential tasks, such as processing language, understanding time-series data, or analyzing events, context and temporal dependencies play a vital role.

- Traditional neural networks process inputs independently without considering the sequence or context.
- RNNs overcome this limitation by introducing recurrent connections that allow information to flow across time steps.
- At each time step, an RNN takes an input, produces an output, and updates its hidden state.
- The hidden state acts as memory, capturing information from previous steps and incorporating it with the current input.
- This enables RNNs to reason sequentially and leverage prior context, making them effective in tasks involving sequential data.
- For example, in movie event classification, an RNN can learn to recognize patterns and dependencies among events to make accurate predictions.

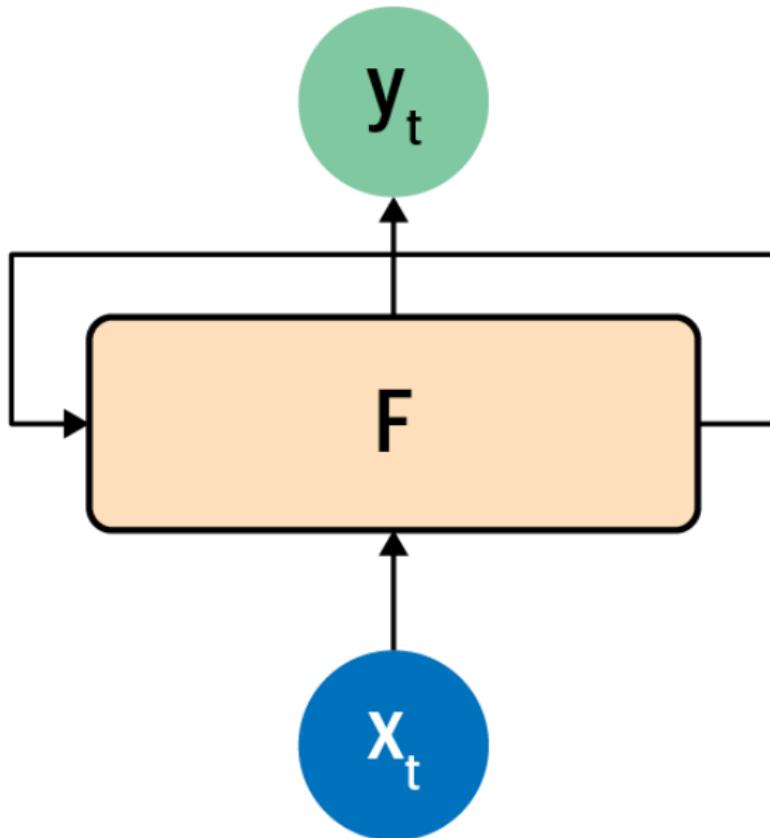


Figure: A recurrent neuron

# Recurrent Neural Networks (RNNs)

- RNNs process sequential data, such as time series or text.
- A value  $x_t$  is fed into the function  $F$  at a time step  $t$ .
- The output  $y_t$  is generated based on  $x_t$ .
- RNNs have a feedback loop, allowing information to flow from one step to another.
- This feedback enables capturing temporal dependencies in the data.

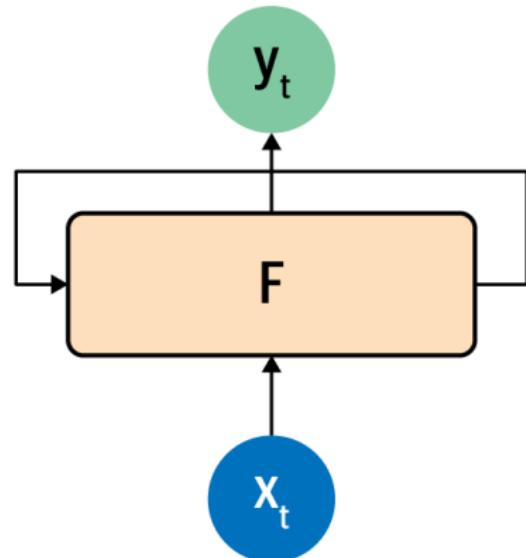


Figure: Recurrent Neural Network (RNN) Architecture

# Recurrent Neural Networks (RNNs)

- RNNs process sequential data.
- At each time step, an RNN takes an input  $x_t$  and produces an output  $y_t$ .
- The RNN maintains a hidden state that captures information from previous steps.
- The hidden state is passed forward to the next time step.

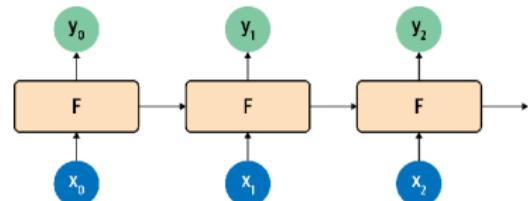


Figure: Recurrent Neural Network (RNN) Architecture

# One-to-One

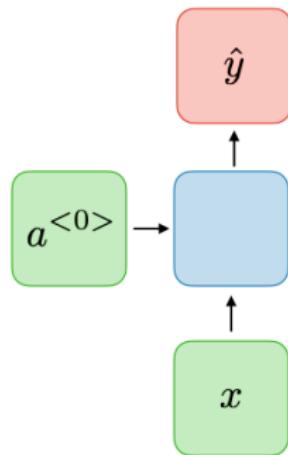


Figure: Illustration of an RNN architecture

## Example: Traditional Neural Network

# One-to-Many

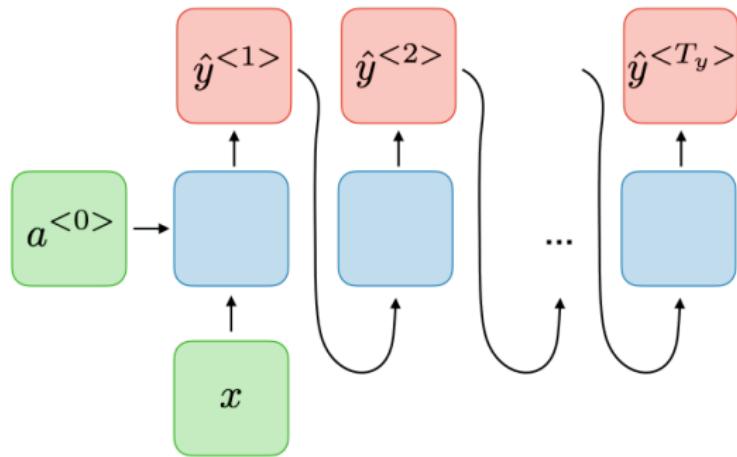


Figure: Illustration of an RNN architecture

## Example: Language Modelling

# Many-to-One

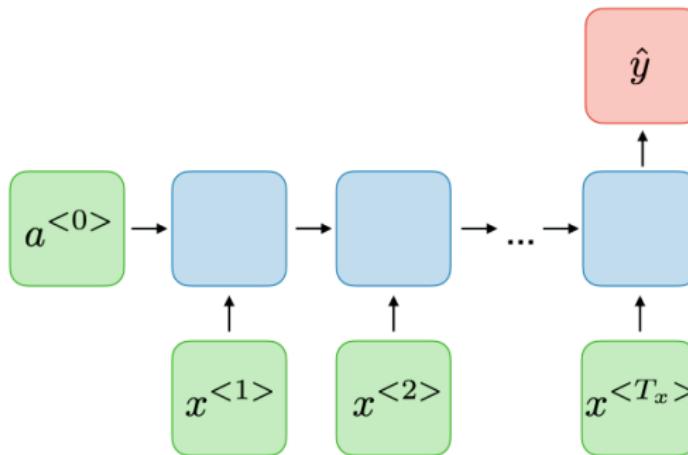


Figure: Illustration of an RNN architecture

**Example:** Sentiment classification

# Many-to-Many

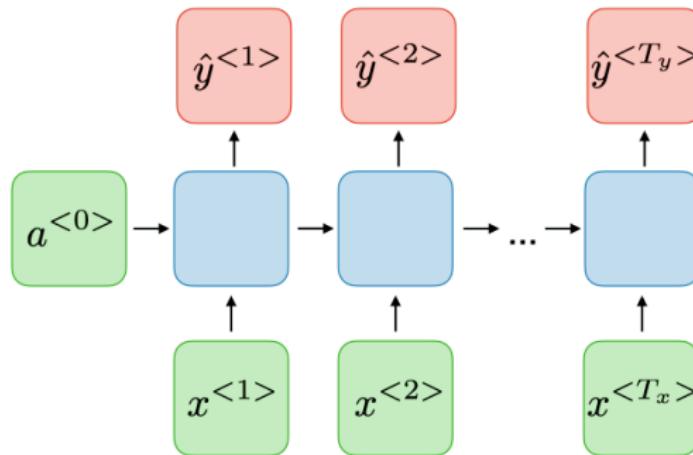


Figure: Illustration of an RNN architecture

**Example:** Name entity recognition

# Recurrent Neural Networks (RNNs)

- RNNs process sequential data.
- At each time step, an RNN takes an input  $x_t$  and produces an output  $y_t$ .
- The RNN maintains a hidden state that captures information from previous steps.
- The hidden state is passed forward to the next time step.

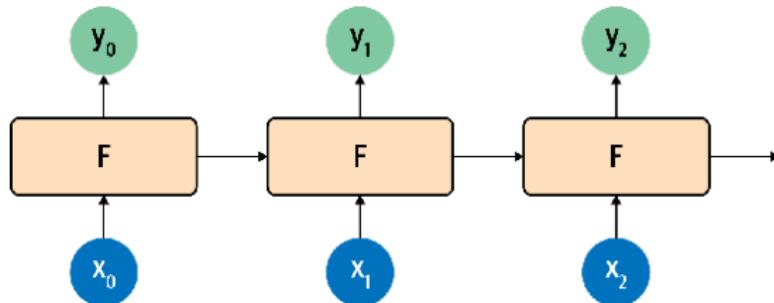


Figure: Recurrent Neural Network (RNN) Architecture

# Why Recurrent Neural Networks?

- Feed-forward neural networks have limitations when it comes to handling sequential data.
- Sequential data, such as time series, text, and speech, cannot be effectively processed by traditional feed-forward networks.
- Feed-forward networks consider only the current input and lack the ability to capture temporal dependencies.
- They do not have memory to retain information about previous inputs.
- RNNs were designed to address these limitations and enable processing of sequential data.
- RNNs can capture dependencies over time and maintain a memory of past inputs.
- This makes RNNs suitable for tasks like speech recognition, language translation, and sentiment analysis.

# Advantages of Recurrent Neural Network (RNN)

- Ability To Handle Variable-Length Sequences
  - RNNs can handle input sequences of variable length.
  - Well-suited for speech recognition, natural language processing, and time series analysis.
- Memory Of Past Inputs
  - RNNs have a memory of past inputs.
  - Captures information about the context of the input sequence.
  - Useful for tasks such as language modeling.
- Parameter Sharing
  - RNNs share the same set of parameters across all time steps.
  - Reduces the number of parameters to be learned.
  - Can lead to better generalization.
- Non-Linear Mapping
  - RNNs use non-linear activation functions.
  - Allows them to learn complex, non-linear mappings between inputs and outputs.

# Advantages of Recurrent Neural Network (RNN) (contd.)

- Sequential Processing
  - RNNs process input sequences sequentially.
  - Computationally efficient and easy to parallelize.
- Flexibility
  - RNNs can be adapted to a wide range of tasks and input types.
  - Suitable for text, speech, and image sequences.
- Improved Accuracy
  - RNNs achieve state-of-the-art performance on various sequence modeling tasks.
  - Language modeling, speech recognition, and machine translation, among others.

# Disadvantages of Recurrent Neural Network (RNN)

- Vanishing And Exploding Gradients
  - RNNs can suffer from the problem of vanishing or exploding gradients during training.
  - It becomes difficult to effectively train the network when gradients become too small or too large.
- Computational Complexity
  - RNNs can be computationally expensive to train, especially with long sequences.
  - Each input needs to be processed sequentially, leading to slow training times.
- Difficulty In Capturing Long-Term Dependencies
  - RNNs may struggle to capture long-term dependencies in the input sequence.
  - Gradients can become very small, causing the network to forget important information.

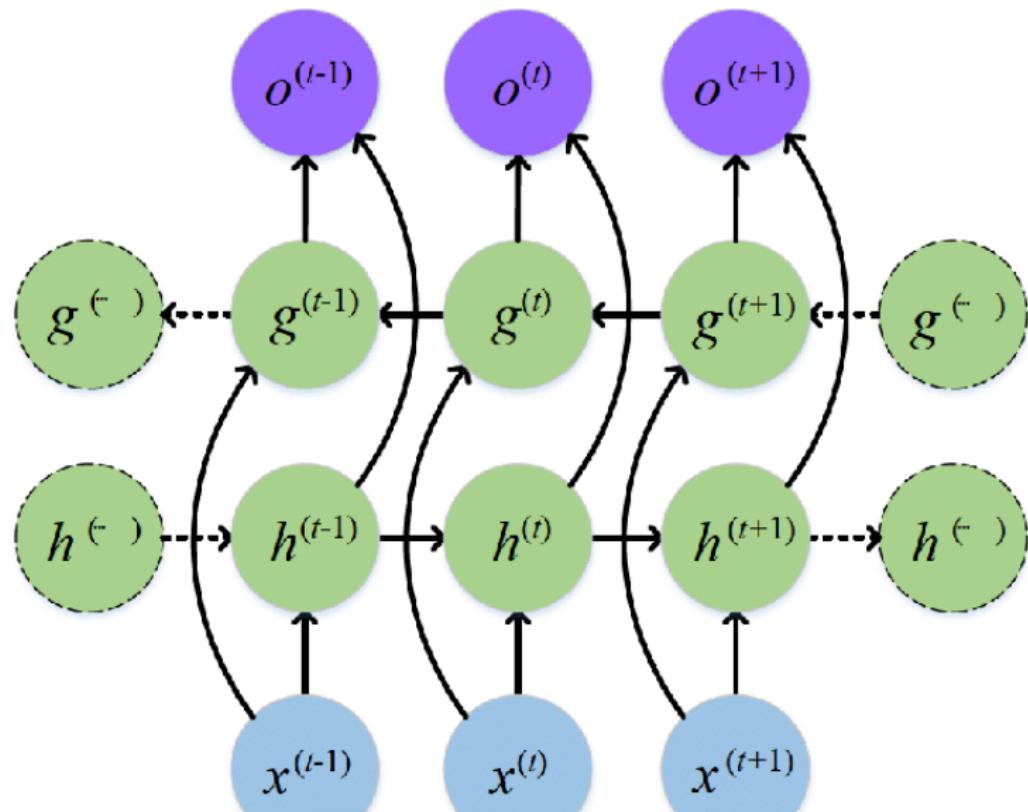
# Disadvantages of Recurrent Neural Network (RNN) (contd.)

- Lack Of Parallelism
  - RNNs are inherently sequential, limiting parallelism in computation.
  - This can impact the speed and scalability of the network.
- Difficulty In Choosing The Right Architecture
  - There are many variants of RNNs with different pros and cons.
  - Selecting the appropriate architecture for a specific task can be challenging.
- Difficulty In Interpreting The Output
  - The output of an RNN can be challenging to interpret, especially for complex inputs.
  - Understanding how the network makes predictions can be difficult.

# Bidirectional Recurrent Neural Networks (RNNs)

- Bidirectional RNNs process input sequences in both forward and backward directions.
- They consist of two separate hidden layers: one processing the sequence in the forward direction and the other in the backward direction.
- Each hidden layer maintains its own hidden state, capturing information from the past and future contexts.
- The outputs of both hidden layers are combined to produce the final prediction or representation.

# Bidirectional Recurrent Neural Networks (RNNs)



# Introduction to Long Short-Term Memory (LSTM) Networks

VVIT@2025

November 8, 2025

# The Problem of Long-Term Dependencies

In certain tasks, recent information is sufficient for the present task. RNNs can effectively use past information when the gap between relevant information and its application is small.

- : In certain tasks, recent information is sufficient for the present task.
- : RNNs can effectively use past information when the gap between relevant information and its application is small.
- Language Modeling Example:
  - Language models predict the next word based on previous words.
  - Example: In the phrase "the sun rises in the," predicting the next word is straightforward.
  - The context of "the sun rises in" is sufficient to conclude that the next word will be "morning."
  - RNNs can successfully utilize the recent context to make accurate predictions.

# Short term dependancies in RNN

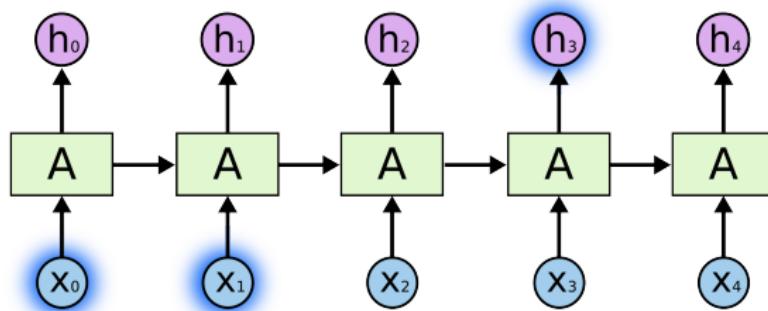


Figure: Example of a Recurrent Neural Network

# Long term dependancies in RNN

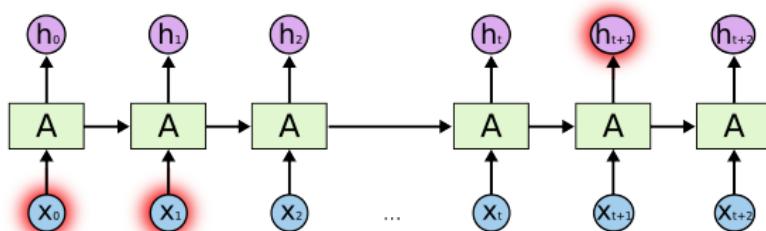


Figure: Long term dependancies in RNN

- I Grew up in Delhi. I did my schooling and college studies there. I speak fluent Hindi.

# Long Short-Term Memory Networks (LSTMs)

In standard RNNs, the repeating module will have a very simple structure, such as a single tanh layer.

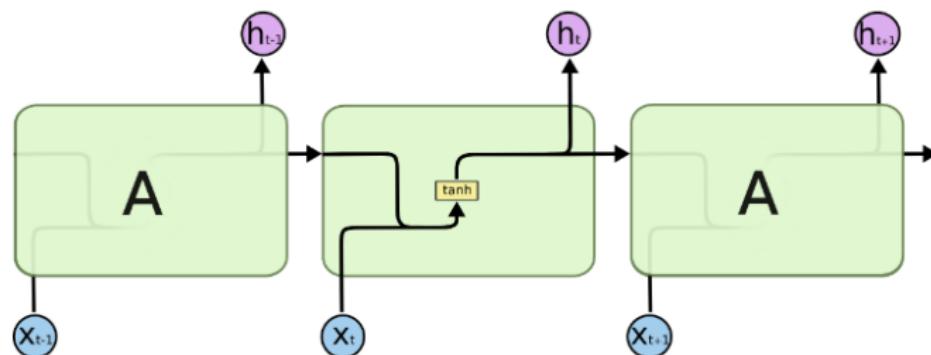


Figure: Long Short-Term Memory (LSTM) Structure

# Long Short-Term Memory (LSTM)

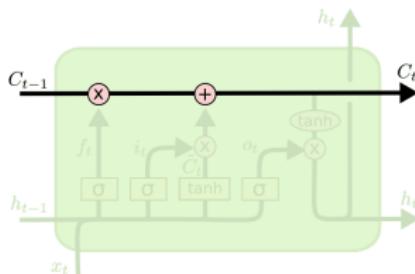


Figure: LSTM Cell Diagram

- The key to LSTMs is the cell state, represented by the horizontal line running through the top of the diagram.
- The cell state acts like a conveyor belt, running straight down the entire chain with minor linear interactions.
- Information can flow along the cell state largely unchanged.
- LSTMs have gates that regulate the removal or addition of information to the cell state.

# What is LSTM?

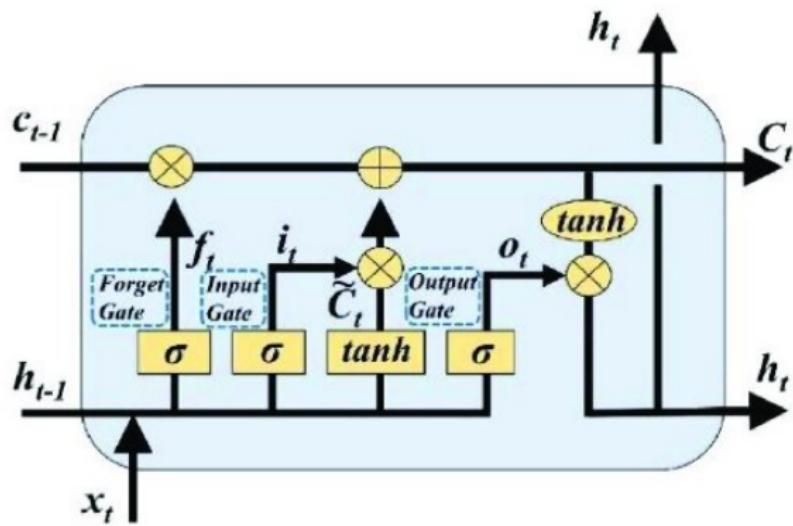
- LSTM stands for Long Short-Term Memory.
- A type of Recurrent Neural Network (RNN) that can learn long-term dependencies.
- Developed to overcome the vanishing gradient problem faced by traditional RNNs.
- Effective for tasks involving sequence prediction, language modeling, and more.

# Why Use LSTM?

- Capable of learning from long sequences.
- Retains important information for longer periods.
- Controls the flow of information using gates.
- Widely used in speech recognition, machine translation, and time-series forecasting.

# LSTM Architecture Overview

- LSTMs have a special architecture designed to remember and forget information.
- Composed of memory cells and gates (input, forget, output).
- Gates regulate the flow of information into and out of the cell state.



# Forget Gate

## Purpose

The forget gate decides what information from the cell state should be discarded or kept.

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) \quad (1)$$

- $f_t$  is the forget gate output.
- $\sigma$  is the sigmoid function.
- $W_f$  is the weight matrix for the forget gate.
- $h_{t-1}$ : Previous hidden state.
- $x_t$ : Current input.
- $b_f$ : Bias term.

# Input Gate

## Purpose

The input gate decides which new information is stored in the cell state.

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \quad (2)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \quad (3)$$

- $i_t$  is the input gate output.
- $\tilde{C}_t$  is the candidate cell state.
- $\tanh$ : Hyperbolic tangent function.
- $W_i$  and  $W_C$  are weight matrices.
- $b_i$  and  $b_C$  are bias terms.

# Output Gate

## Purpose

The output gate decides what part of the cell state should be output to the next time step.

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) \quad (4)$$

$$h_t = o_t \cdot \tanh(C_t) \quad (5)$$

- $o_t$  is the output gate output.
- $h_t$  is the new hidden state.
- $W_o$  is the weight matrix for the output gate.
- $b_o$  is the bias term.

# Updating Cell State

## Cell State Equation

$$C_t = f_t \cdot C_{t-1} + i_t \cdot \tilde{C}_t$$

- $C_t$ : Current cell state.
- $C_{t-1}$ : Previous cell state.
- $f_t \cdot C_{t-1}$ : Forget gate's decision to discard or keep previous information.
- $i_t \cdot \tilde{C}_t$ : Input gate's decision to add new information.

# Applications of LSTM

- Language Modeling: Predicting the next word in a sentence.
- Machine Translation: Translating text from one language to another.
- Speech Recognition: Converting spoken language into text.
- Time Series Forecasting: Predicting future values based on past data.

# GRU cell in practice

- Simplified LSTM

- No cell state
- Two gates (instead of three)
- Fewer weights

- Update equations

$$\text{Reset gate: } r_t = \sigma(W^{(ir)}\bar{x}_t + W^{(hr)}h_{t-1})$$

$$\text{Update gate: } z_t = \sigma(W^{(iz)}\bar{x}_t + W^{(hz)}h_{t-1})$$

$$\text{Process input: } \tilde{h}_t = \tanh(W^{(i\tilde{h})}\bar{x}_t + r_t * (W^{(h\tilde{h})}h_{t-1}))$$

$$\text{Hidden state update: } h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

$$\text{Output: } y_t = h_t$$

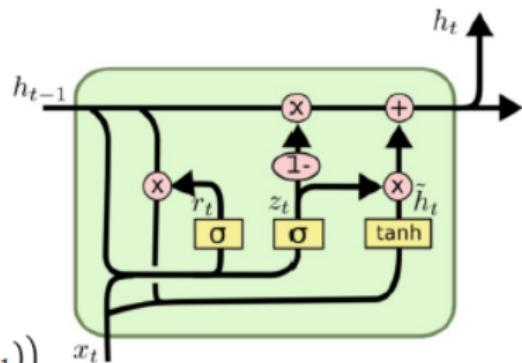


Figure: GRU Cell